

In this overview and graph the results of the CUDA experiments are compared with the results from previous experiments 1.1, 2.1 and 3.1. In each table the 'winner' from each previous experiment is selected to avoid an unnecessarily big table.

We can also see the time does not increase that much when adding more points. Probably this is because (the high parallelism in the graphics card) the time needed to transfer the data is responsible for the most time.

However, from  $10^4$  points the multithread implementations start to catch up and even outpace the sequential implementations. CUDA is faster than every other implementation.

The more points we add the higher the speedup gets.

